What exactly has TabPFN learned to do?

Calvin McCarter

More results here!



Background on TabPFN

TabPFN^[1] is a prior-data fitted network (PFN) that was pretrained to perform in-context learning on fresh tabular classification problems.

- Meta-learned Transformer model (26M parameters)
- Trained using synthetically-generating data: structural causal models are randomly generated, then training datasets are sampled from each causal model.
- On fresh classification tasks, no training (i.e. weight updating) is needed; instead, training data is given as context.

[1] TabPFN: A Transformer That Solves Small Tabular Classification Problems in a Second. N. Hollmann, S. Muller, K. Eggensperger, F. Hutter. ICLR 2023.

Questions

- Was TabPFN overfit to tabular benchmarks?
- Has TabPFN learned fundamental principles of statistical learning?

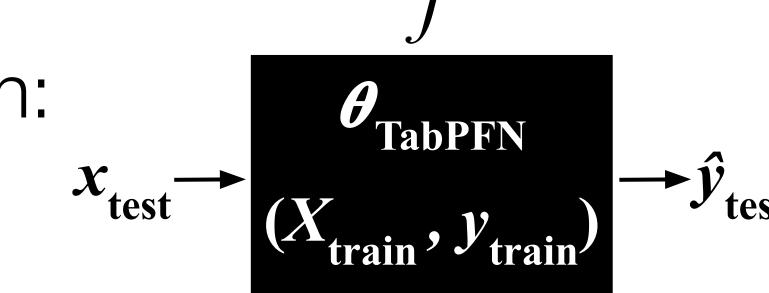
Our approach

Understand what TabPFN-the-model-artifact does:

Observe TabPFN's test outputs as a function of its test inputs, for a variety of training sets:

- Simple settings, so we can plot the learned functions
- Out-of-domain settings, so we can evaluate generalization

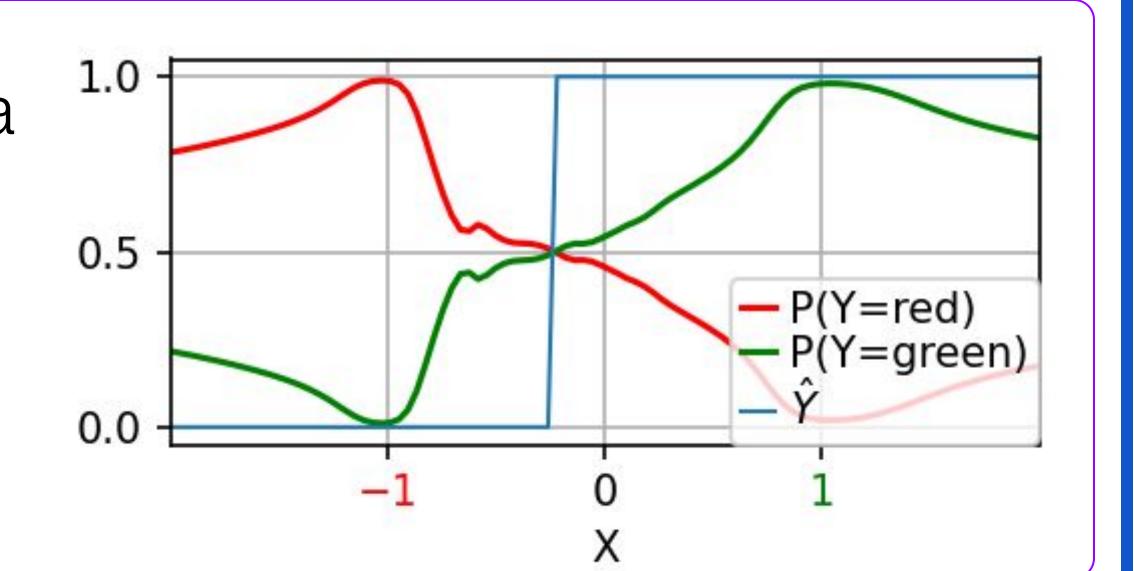
TabPFN, with training data, is just a function:



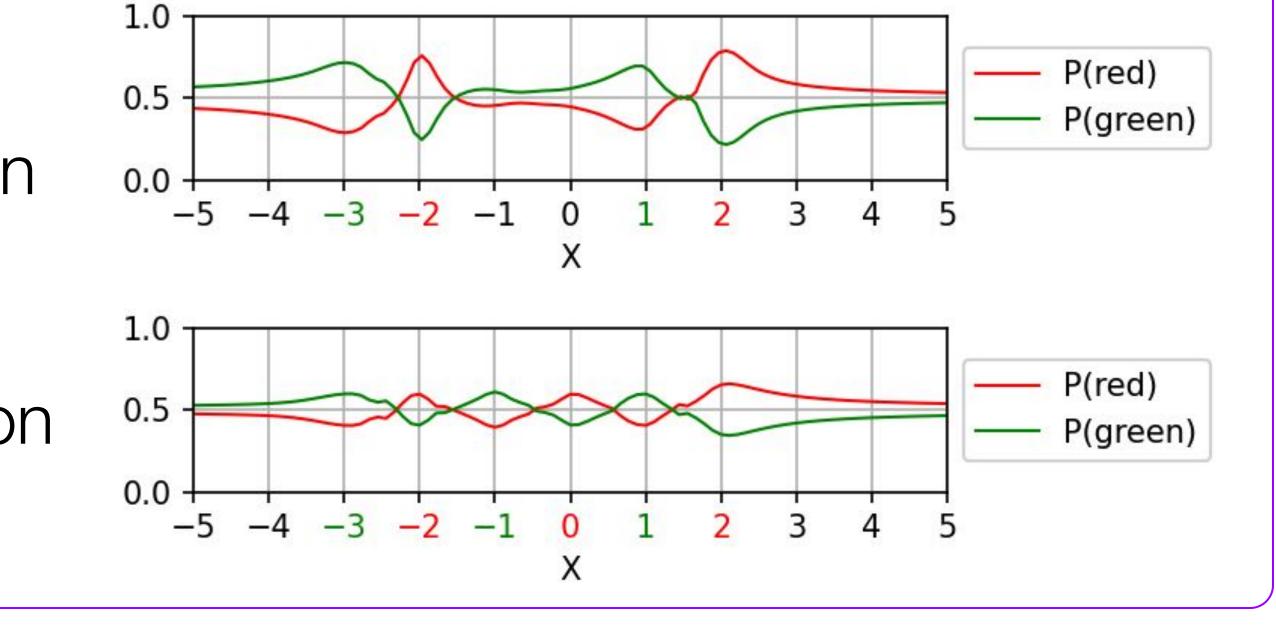
Simple synthetic settings

1d binary classification

- → expects right-skewed data
- → has epistemic uncertainty outside the range of its training data



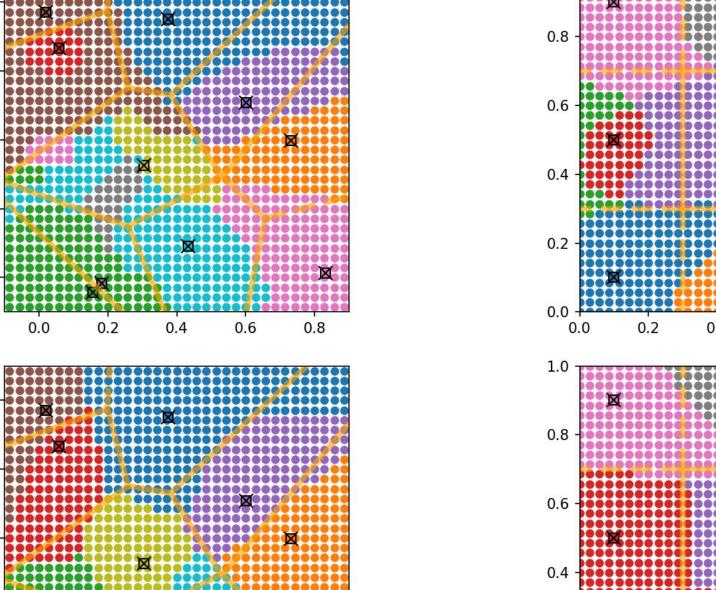
- → cannot perform periodic interpolation
- → cannot perform periodic extrapolation



2d multi-class classification

Has TabPFN learned the principle of nearest-neighbor classification? Not without ensembling!

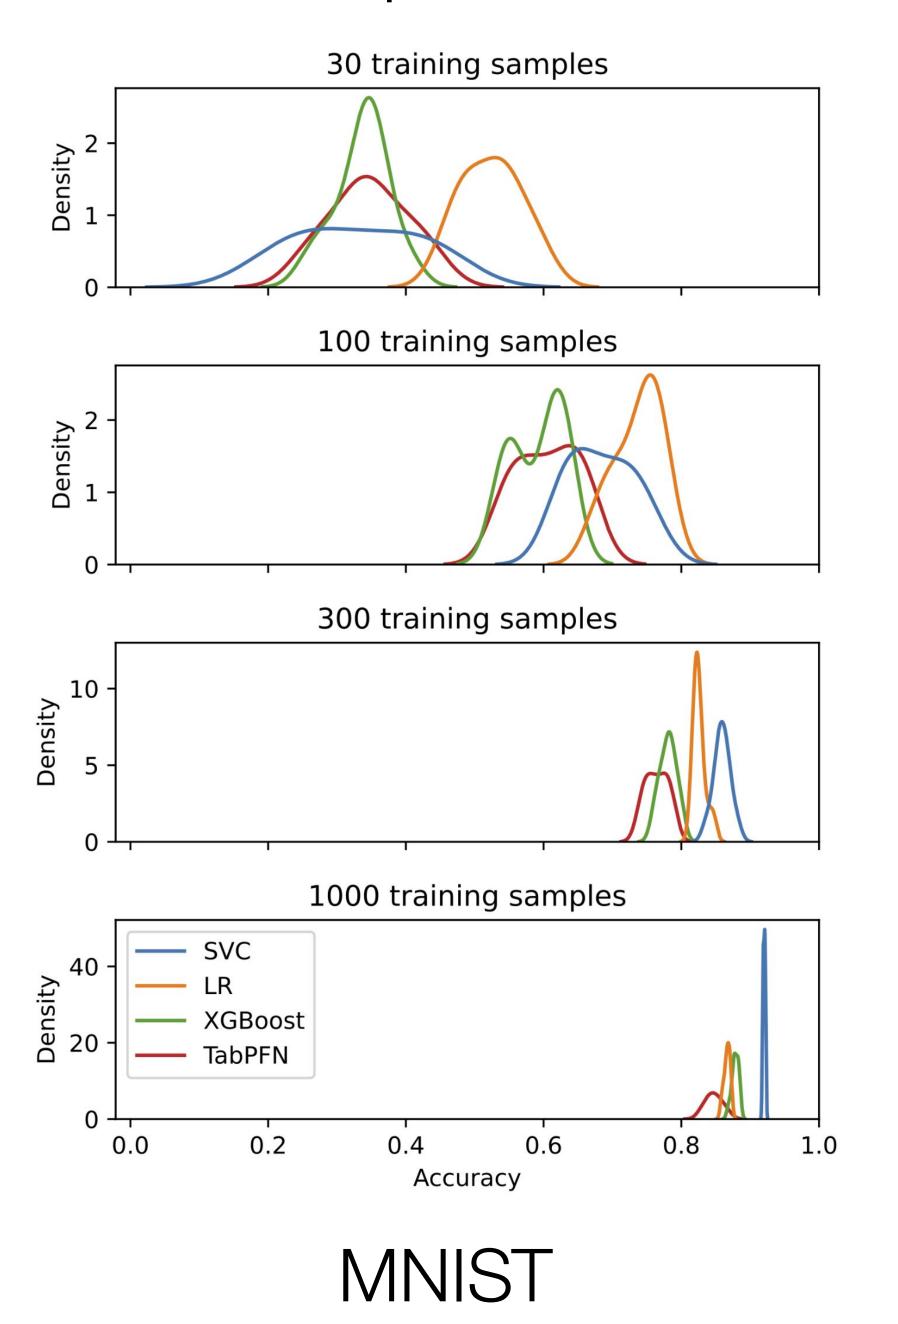


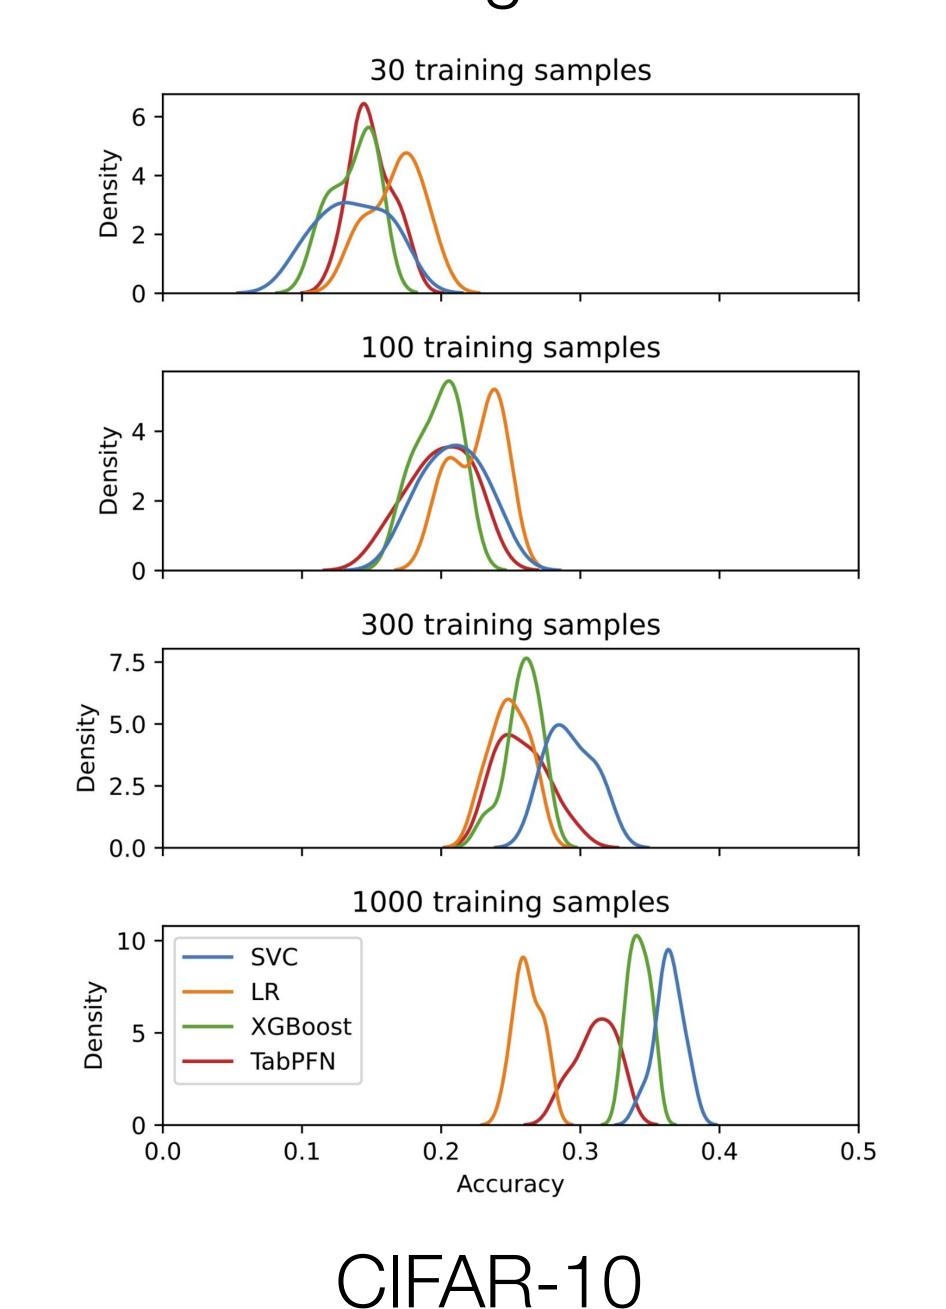




Out-of-domain settings

Compare support vector (SVC), logistic regression (LR), XGBoost, and TabPFN on problems in computer vision and genomics.





Conclusions

TabPFN:

- o does not appear overfit to tabular benchmarks.
- o appears to sometimes implement fundamental statistical learning principles.

Pretrained tabular ML models should be evaluated like LLMs:

- Theoretical guarantees and benchmark results are not enough!
- Evaluate specific model artifacts, not only methods.